

Object Oriented Programming - Exercise 4: AVL Trees

Contents

| | | |
|----------|---------------------------------|----------|
| 1 | Goals | 1 |
| 2 | Submission Details | 2 |
| 3 | Introduction | 2 |
| 4 | API | 2 |
| 4.1 | Package | 2 |
| 4.2 | Methods | 2 |
| 4.3 | Constructors | 3 |
| 4.4 | Static Methods | 4 |
| 4.5 | Comments | 4 |
| 5 | Analyzing the AVL Tree | 4 |
| 6 | Submission Requirements | 5 |
| 6.1 | README | 5 |
| 6.2 | Automatic Testing | 5 |
| 6.3 | Submission Guidelines | 6 |
| 7 | Technical Issues | 6 |
| A | Test Input File Format | 7 |
| A.1 | Examples | 7 |
| A.1.1 | | 7 |
| A.1.2 | | 8 |
| A.2 | Some Clarifications | 8 |

1 Goals

1. Implementing a data structure you have learned about in the Data Structures course - the AVL Tree.
2. Experimenting with this data structure.

2 Submission Details

- Submission Deadline: **Sunday, 21/5/2015, 23:55**
- This exercise will be done alone (**not in pairs**).
- You may use the `java.util.List` interface and two of its implementations – `java.util.LinkedList` and `java.util.ArrayList`. You may also use the `java.lang.Iterable<T>` interface and the `java.util.Iterator` class, the `Math` class and any `Exception` classes you want. You **may not** use any other class that you didn't write yourself.

3 Introduction

An AVL tree is a self-balancing binary search tree. AVL trees maintain the *AVL property* - for each node, the difference between the heights of both of its sub-trees is at most 1.¹ This is done by re-balancing the tree after each insertion and deletion operation. You've recently learned about the theoretical background of AVL trees in the Data Structures (DaSt) course.²

In this exercise, you will implement an AVL tree of integers based on the pseudo-code shown in the DaSt lecture and `tirgul`.

4 API

You are required to submit a class named `AvlTree` that implements `Iterable<Integer>` and the following API.

You are allowed (and encouraged!) to write and submit more class(es) as you see fit. You can also add methods of your own but make sure that you don't change the supplied API; you are only allowed to add methods with the `private` and the `default-package` access modifier.

4.1 Package

Both `AvlTree` and any other class you implement in this exercise should be placed in the `oop.ex4.data_structures` package.

4.2 Methods

- ```
/**
 * Add a new node with the given key to the tree.
 *
 * @param newValue the value of the new node to add.
 * @return true if the value to add is not already in the tree and it was successfully added,
 * false otherwise.
 */
public boolean add(int newValue);
```

---

<sup>1</sup>The height of a tree is defined as the length of the longest downward path from the root to any of the leaves.

<sup>2</sup>See <http://moodle.cs.huji.ac.il/cs13/file.php/67109/tirgul8.pdf>.

- ```
/**
 * Check whether the tree contains the given input value.
 *
 * @param val the value to search for.
 * @return the depth of the node (0 for the root) with the given value if it was found in
 * the tree, -1 otherwise.
 */
public int contains(int searchVal);
```
- ```
/**
 * Removes the node with the given value from the tree, if it exists.
 *
 * @param toDelete the value to remove from the tree.
 * @return true if the given value was found and deleted, false otherwise.
 */
public boolean delete(int toDelete);
```
- ```
/**
 * @return the number of nodes in the tree.
 */
public int size();
```
- ```
/**
 * @return an iterator on the Avl Tree. The returned iterator iterates over the tree nodes
 * in an ascending order, and does NOT implement the remove() method.
 */
public Iterator<Integer> iterator();
```

## 4.3 Constructors

- ```
/**
 * The default constructor.
 */
public AvlTree();
```
- ```
/**
 * A constructor that builds the tree by adding the elements in the input array one by
 * one. If a value appears more than once in the list, only the first appearance is added.
 *
 * @param data the values to add to tree.
 */
public AvlTree(int[] data);
```
- ```
/**
 * A copy constructor that creates a deep copy of the given AvlTree. This means that for
 * every node or any other internal object of the given tree, a new, identical object, is
```

```

    * instantiated for the new tree (the internal object is not simply referenced from it). The
    * new tree must contain all the values of the given tree, but not necessarily in the same
    * structure.
    *
    * @param avlTree an AVL tree.
    */
    public AvlTree(AvlTree avlTree);

```

4.4 Static Methods

```

• /**
   * Calculates the minimum number of nodes in an AVL tree of height h.
   *
   * @param h the height of the tree (a non-negative number) in question.
   * @return the minimum number of nodes in an AVL tree of the given height.
   */
    public static int findMinNodes(int h);

```

For example, `findMinNodes(3)`, as shown in Figure 1, should return 7

4.5 Comments

1. The `add(int newValue)` method will return false and leave the tree unchanged if `newValue` already exists in it.
2. The data constructor can receive any array of int primitives, which means **you can't assume the array is sorted**. You can't assume you won't get a null reference.
3. **Do not** paste code copied from the PDF file. It contains unwanted characters that will fail compilation.
4. A tree with only a single node is of height zero, or has $h = 0$. A tree with a root that has one or two sons (each of which has no sons) is of height one, or $h = 1$. Etc.
5. When writing a class implementing an interface only partially - opting not to implement some optional methods of the interface - you can throw an `UnsupportedOperationException` from the methods you wish to leave unimplemented. Can you figure out how throwing this specific exception does not violate the interface's API and method signature?

5 Analyzing the AVL Tree

In Figure 1 you can see an AVL tree of height 3. This tree may seem unbalanced, but is in fact a valid AVL tree resulted from a specific insertion order. Notice that this example shows a tree of height 3 with the minimal number of nodes (if you delete any node, the height of the tree will become 2). Find a series of 12 numbers, such that when they are inserted into an empty AVL tree one-by-one, the result will be a tree of height 4 (insertions only, no deletions). Write a full answer in the README file.

Hint: We suggest you start by figuring out the order in which the nodes were inserted into the example tree, and then continue to create the larger one.




Figure 1: an AVL tree of height 3 with the minimal number of nodes

6 Submission Requirements

6.1 README

Please address the following points in your README file:

1. Describe which class(es) (if any) you wrote as part of your implementation of an AVL tree, other than `AvlTree`. The description should include the purpose of each class, its important methods and its interaction with the `AvlTree` class.
2. Describe your implementation of the methods `add()` and `delete()`. The description should include the general workflow in each of these methods. You should also indicate which helper methods you implemented for each of them, and which of these helper methods are shared by both of them (if any).
3. Your README file should also answer the question from Section 5.

6.2 Automatic Testing

Our automatic tests will do the following:

1. Make sure your code compiles and contains a README file and a QUESTIONS file. This includes verifying that your `AvlTree` class follows the API defined in Section 4 (including the correct package names).
2. Test each of your methods on simple inputs, in order to test their basic functionality.
3. Test your tree on the four interesting AVL scenarios (i.e., the four rotation cases).
4. Test your code on more complicated scenarios that include large inputs and multiple add/delete operations.

5. As always, we will run unseen tests which weigh 10% of the automatic testers.

We will compile your class with our own tester and run a set of tests. Each such test is defined by an input file, which creates an AVL tree and performs a set of operations on it. The format of the input files is described in Appendix A.

6.3 Submission Guidelines

You should submit a file named `ex4.jar` containing all the `.java` files of your program, as well as the `README` and `QUESTIONS` file. Please note the following:

- Files should be submitted in the original directory hierarchy of their packages.
- No `.class` files should be submitted.
- There is no need to submit any testers.
- Your program must compile without any errors or warnings. All the warnings we will look for are also reported by Eclipse. If you want to check for them using the console, you can use the following customized compilation command: `javac Xlint: rawtypes Xlint: static Xlint: empty Xlint: divzero Xlint: deprecation file1.java file2.java` Note: Copy-pasting the command line above to the shell might result in some invalid characters.
- `javadoc` should compile without any errors or warnings.

You may use the following unix command to create the jar files:

```
jar -cvf ex4.jar README QUESTIONS oop/ex4/data_structures/*.java
```

This command should be run from the main project directory (that is, where the `oop` directory is found).

7 Technical Issues

A file named `ex4_files.zip` can be found in the course website. This zip contains the test files on which your program will be tested. When you submit your exercise, we will run your code against all the input files, and you will receive a response file with the tests you failed on (if any). Error messages will contain some information about the problem. For example, you may receive the following error:

```
runTests[17] (Ex4Tester):  ex4/tests/020.txt (# simple right-left rotation after delete):  
line number 7:  expected:<true> but was:<false>
```

This means that your program failed the test in the file `020.txt`, when trying to execute the delete operation. More specifically, line 7 of the test failed at this stage because the return value of calling `delete()` was `true` instead of `false`.

Good Luck!

Appendices

A Test Input File Format

Each of the automatic tests is incorporated in a file of the following format:

- The first line starts with a ‘#’ sign, and is a comment describing the test (i.e., this line is ignored).
- The second line creates an `AvlTree` object (denoted `myTree`) using one of the constructors defined in Section 4.3:
 - An empty line calls the default constructor.
 - A list of integers (e.g., `1 5 7 12 4`) calls the data constructor with this list of integers.
- The following lines specify a call to one of the methods of the created tree, and are executed sequentially.

They can be one of the following:

- *add number return_value*
call `myTree.add(number)`. Verify that the method’s return value matches *return_value* (either `true` or `false`).
- *delete number return_value*
call `myTree.delete(number)`. Verify that the method’s return value matches *return_value* (either `true` or `false`).
- *contains number return_value*
call `myTree.contains(number)`. Verify that the method’s return value matches *return_value* (a non-negative integer or -1).
- *size return_value*
call `myTree.size()`. Verify that the method’s return value matches *return_value* (a non-negative integer).
- *copy*
call `myTree2 = newAvlTree(myTree)`. Verify that the copy constructor functions correctly by iterating over `myTree` and `myTree2` and verifying they contain the same values.
- *minNode number return_value*
call `AvlTree.findMinNodes(number)`. Verify that the method’s return value matches *return_value*.

A.1 Examples

A.1.1

Consider the following example test file:

```
# a simple creation of a tree and adding the number 5
```

```
add 5 true  
size 1
```

In this example, the first line is a comment describing the test. The second line is empty, which means calling the empty constructor. The third line calls `myTree.add(5)`, and then verifies that the command succeeded (by verifying that the return value of the `add` method is `true`). The last line verifies that the tree size is now 1 (after adding a single element).

A.1.2

Consider another example:

```
# calling the data constructor with the values (1 2 3) , deleting 3, and then searching for 4  
1 2 3  
delete 3 true  
contains 4 -1
```

The first line is, again, a comment describing the test. The second line is a list of numbers, which means calling the data constructor with the integer array (1 2 3). The third line calls `myTree.delete(3)`, and then verifies that the command succeeded. The last line calls `myTree.contains(4)` and verifies that 4 is not found in the tree (return value should be -1).

A.2 Some Clarifications

Please note the following:

- Your code passes a given test only if your methods return the correct values in **each** of the method calls.
- A test file may include numerous calls to `add`, `delete`, `contains` and `size`.